



Positions: learned vs. NoPE vs. RoPE
Training loss



Validation loss (from iteration 500)



Same default training settings and seed 1337; one architecture change per run.
Losses: evaluation mode, 200 batches per split. Dots highlight best validation checkpoints.