



LayerNorm vs. RMSNorm: default Shakespeare training

Training loss



Validation loss (from iteration 500)



Same configuration and seed (1337); all 13 normalization layers replaced; epsilon = 1e-5.
Both losses use evaluation mode and 200 batches per split. Markers highlight the best validation loss.